

Can I Run AI locally?

Find out which AI models your machine can actually run.

RTX 4090 Laptop (16 GB) 4 GB * ~576 GB/s 32 GB 9,728
Estimates based on browser APIs. Actual specs may vary.

33 RUNS GREAT | **6 RUNS WELL** | **1 DECENT** | **3 TIGHT FIT** | **20 BARELY RUNS** | **14 TOO HEAVY**



















Q Search models... / All grades All tasks All providers All licenses Sort: Score

List

Llama 3.1 8B		1 year ago	4.6 GB	128K ctx	~88 tok/s	RUNS GREAT 92/100
Qwen 3.5 9B		2mo ago	5.1 GB	32K ctx	~79 tok/s	RUNS GREAT 91/100
Phi-4 14B		1 year ago	7.7 GB	16K ctx	~52 tok/s	RUNS WELL 77/100
GPT-OSS 20B		8mo ago	11.3 GB	128K ctx	~36 tok/s	DECENT 61/100
Mistral Small 3.1 24B		1 year ago	12.8 GB	128K ctx	~31 tok/s	TIGHT FIT 54/100
Gemma 3 27B		1 year ago	14.3 GB	128K ctx	~28 tok/s	BARELY RUNS 37/100
Qwen 2.5 Coder 32B		1 year ago	16.9 GB	128K ctx	~13 tok/s	BARELY RUNS 24/100
Qwen 3 32B		1 year ago	16.9 GB	128K ctx	~13 tok/s	BARELY RUNS 24/100
DeepSeek R1 Distill 32B		1 year ago	16.9 GB	64K ctx	~13 tok/s	BARELY RUNS 24/100
Llama 3.3 70B		1 year ago	36.4 GB RAM	128K ctx	~1 tok/s	BARELY RUNS 8/100
Llama 4 Scout 17B		1 year ago	56.3 GB	128K ctx	0 tok/s	TOO HEAVY 0/100
GPT-OSS 120B		8mo ago	60.4 GB	128K ctx	0 tok/s	TOO HEAVY 0/100
Devstral 2 123B		4mo ago	63.5 GB	256K ctx	0 tok/s	TOO HEAVY 0/100
DeepSeek R1		1 year ago	344.2 GB	64K ctx	0 tok/s	TOO HEAVY 0/100
DeepSeek V3.2		4mo ago	351.4 GB	128K ctx	0 tok/s	TOO HEAVY 0/100
Kimi K2		9mo ago	512.7 GB	128K ctx	0 tok/s	TOO HEAVY 0/100

All models

Phi-3.5 Mini		1 year ago	2.4 GB	128K ctx	~168 tok/s	RUNS GREAT 95/100
Phi-4 Mini Reasoning		1 year ago	2.4 GB	16K ctx	~168 tok/s	RUNS GREAT 95/100
Qwen 3 4B		1 year ago	2.5 GB	32K ctx	~161 tok/s	RUNS GREAT 95/100
Gemma 3 4B		1 year ago	2.5 GB	128K ctx	~161 tok/s	RUNS GREAT 95/100

Qwen 3.5 4B	 	2mo ago	2.5 GB	16%	32K ctx	~161 tok/s	RUNS GREAT	95/100
Gemma 4 E2B IT	 	5d ago	3.1 GB	19%	256K ctx	~130 tok/s	RUNS GREAT	95/100
Gemma 4 E2B	 	5d ago	3.1 GB	19%	256K ctx	~130 tok/s	RUNS GREAT	95/100
Llama 3.2 3B	 	1 year ago	2 GB	13%	128K ctx	~202 tok/s	RUNS GREAT	94/100
SmolLM3 3B		9mo ago	2 GB	13%	128K ctx	~202 tok/s	RUNS GREAT	94/100
Qwen 2.5 Coder 1.5B		1 year ago	1.3 GB	8%	32K ctx	~310 tok/s	RUNS GREAT	93/100
DeepSeek R1 1.5B	 	1 year ago	1.3 GB	8%	64K ctx	~310 tok/s	RUNS GREAT	93/100
Qwen 3 1.7B	 	1 year ago	1.4 GB	9%	32K ctx	~288 tok/s	RUNS GREAT	93/100
Qwen 3.5 2B	 	2mo ago	1.5 GB	9%	32K ctx	~269 tok/s	RUNS GREAT	93/100
Gemma 2 2B		1 year ago	1.5 GB	9%	8K ctx	~269 tok/s	RUNS GREAT	93/100
Mistral 7B v0.3	 	1 year ago	4.1 GB	26%	32K ctx	~98 tok/s	RUNS GREAT	93/100
Qwen 2.5 7B	 	1 year ago	4.1 GB	26%	128K ctx	~98 tok/s	RUNS GREAT	93/100
Qwen 2.5 Coder 7B	 	1 year ago	4.1 GB	26%	128K ctx	~98 tok/s	RUNS GREAT	93/100
DeepSeek R1 Distill 7B	 	1 year ago	4.1 GB	26%	64K ctx	~98 tok/s	RUNS GREAT	93/100
Qwen 3.5 0.8B	 	2mo ago	0.9 GB	6%	32K ctx	~448 tok/s	RUNS GREAT	92/100
Llama 3.2 1B	 	1 year ago	1 GB	6%	128K ctx	~403 tok/s	RUNS GREAT	92/100
Gemma 3 1B		1 year ago	1 GB	6%	32K ctx	~403 tok/s	RUNS GREAT	92/100
TinyLlama 1.1B		2y ago	1.1 GB	7%	2K ctx	~367 tok/s	RUNS GREAT	92/100
Gemma 4 E4B IT	 	5d ago	4.6 GB	29%	256K ctx	~88 tok/s	RUNS GREAT	92/100
Gemma 4 E4B	 	5d ago	4.6 GB	29%	256K ctx	~88 tok/s	RUNS GREAT	92/100
Qwen 3 8B	  	1 year ago	4.6 GB	29%	128K ctx	~88 tok/s	RUNS GREAT	92/100
Ministral 8B	 	1 year ago	4.6 GB	29%	32K ctx	~88 tok/s	RUNS GREAT	92/100
Qwen 3 0.6B	 	1 year ago	0.8 GB	5%	32K ctx	~504 tok/s	RUNS GREAT	91/100
Gemma 2 9B		1 year ago	5.1 GB	32%	8K ctx	~79 tok/s	RUNS GREAT	91/100
GLM-4 9B	 	1 year ago	5.1 GB	32%	128K ctx	~79 tok/s	RUNS GREAT	91/100
Nemotron Nano 9B v2	 	10mo ago	5.1 GB	32%	128K ctx	~79 tok/s	RUNS GREAT	91/100
Llama 3.2 11B Vision	  	1 year ago	6.1 GB	38%	128K ctx	~66 tok/s	RUNS GREAT	85/100
Gemma 3 12B	 	1 year ago	6.6 GB	41%	128K ctx	~61 tok/s	RUNS WELL	83/100
Mistral Nemo 12B	 	1 year ago	6.6 GB	41%	128K ctx	~61 tok/s	RUNS WELL	83/100
Qwen 2.5 14B	 	1 year ago	7.7 GB	48%	128K ctx	~52 tok/s	RUNS WELL	77/100
Qwen 3 14B	  	1 year ago	7.7 GB	48%	128K ctx	~52 tok/s	RUNS WELL	77/100
DeepSeek R1 Distill 14B	 	1 year ago	7.7 GB	48%	64K ctx	~52 tok/s	RUNS WELL	77/100

LFM2 24B		5mo ago	12.8 GB	80%	32K ctx	~31 tok/s	TIGHT FIT	54/100
Devstral Small 2 24B		4mo ago	12.8 GB	80%	256K ctx	~31 tok/s	TIGHT FIT	54/100
Gemma 2 27B		1 year ago	14.3 GB	89%	8K ctx	~28 tok/s	BARELY RUNS	37/100
Gemma 4 26B-A4B IT		5d ago	14.3 GB	89%	256K ctx	~28 tok/s	BARELY RUNS	37/100
Gemma 4 26B-A4B		5d ago	14.3 GB	89%	256K ctx	~28 tok/s	BARELY RUNS	37/100
Qwen 3.5 27B		2mo ago	14.7 GB	92%	256K ctx	~27 tok/s	BARELY RUNS	36/100
Qwen 3 30B-A3B		1 year ago	15.9 GB	99%	128K ctx	~25 tok/s	BARELY RUNS	33/100
Nemotron 3 Nano 30B		10mo ago	15.9 GB	99%	1024K ctx	~25 tok/s	BARELY RUNS	33/100
Qwen 2.5 32B		1 year ago	16.9 GB	106%	128K ctx	~13 tok/s	BARELY RUNS	24/100
EXAONE 4.0 32B		9mo ago	16.9 GB	106%	128K ctx	~13 tok/s	BARELY RUNS	24/100
OLMo 2 32B		1 year ago	16.9 GB	106%	4K ctx	~13 tok/s	BARELY RUNS	24/100
Gemma 4 31B IT		5d ago	17.4 GB	109%	256K ctx	~11 tok/s	BARELY RUNS	23/100
Gemma 4 31B		5d ago	17.4 GB	109%	256K ctx	~11 tok/s	BARELY RUNS	23/100
Command R 35B		2y ago	18.4 GB	115% ↗ RAM	128K ctx	~8 tok/s	BARELY RUNS	15/100
Qwen 3.5 35B-A3B		2mo ago	18.4 GB	115% ↗ RAM	256K ctx	~8 tok/s	BARELY RUNS	15/100
Mixtral 8x7B		2y ago	24.6 GB	154% ↗ RAM	32K ctx	~3 tok/s	BARELY RUNS	10/100
Qwen 2.5 72B		1 year ago	37.4 GB	234% ↗ RAM	128K ctx	~1 tok/s	BARELY RUNS	8/100
Qwen 3.5 122B-A10B		2mo ago	63 GB	394%	256K ctx	0 tok/s	TOO HEAVY	0/100
Mixtral 8x22B		2y ago	72.7 GB	454%	64K ctx	0 tok/s	TOO HEAVY	0/100
Qwen 3 235B-A22B		1 year ago	120.9 GB	756%	128K ctx	0 tok/s	TOO HEAVY	0/100
Qwen 3.5 397B-A17B		2mo ago	203.9 GB	1274%	256K ctx	0 tok/s	TOO HEAVY	0/100
Llama 4 Maverick 17B-128E		1 year ago	205.4 GB	1284%	1024K ctx	0 tok/s	TOO HEAVY	0/100
Llama 3.1 405B		1 year ago	208 GB	1300%	128K ctx	0 tok/s	TOO HEAVY	0/100
Qwen 3 Coder 480B		9mo ago	246.4 GB	1540%	256K ctx	0 tok/s	TOO HEAVY	0/100
DeepSeek V3.1		8mo ago	344.2 GB	2151%	128K ctx	0 tok/s	TOO HEAVY	0/100

Built by [midudev](#) for the local AI community ❤️

Data sourced from [llama.cpp](#), [Ollama](#) and [LM Studio](#).

All product names, logos, and brands are property of their respective owners. Apple, NVIDIA, AMD, Intel, Qualcomm, and all AI model names mentioned on this site are trademarks or registered trademarks of their respective holders. This site is not affiliated with or endorsed by any of these companies.

